HIVE CASE STUDY

EMR CLUSTER CREATION

STEP 1: Login to your AWS account & search EMR services. After the EMR home page appears click on **Create cluster** & follow the steps as mentioned. We have chosen cluster release version **5.29.0** in our case study.

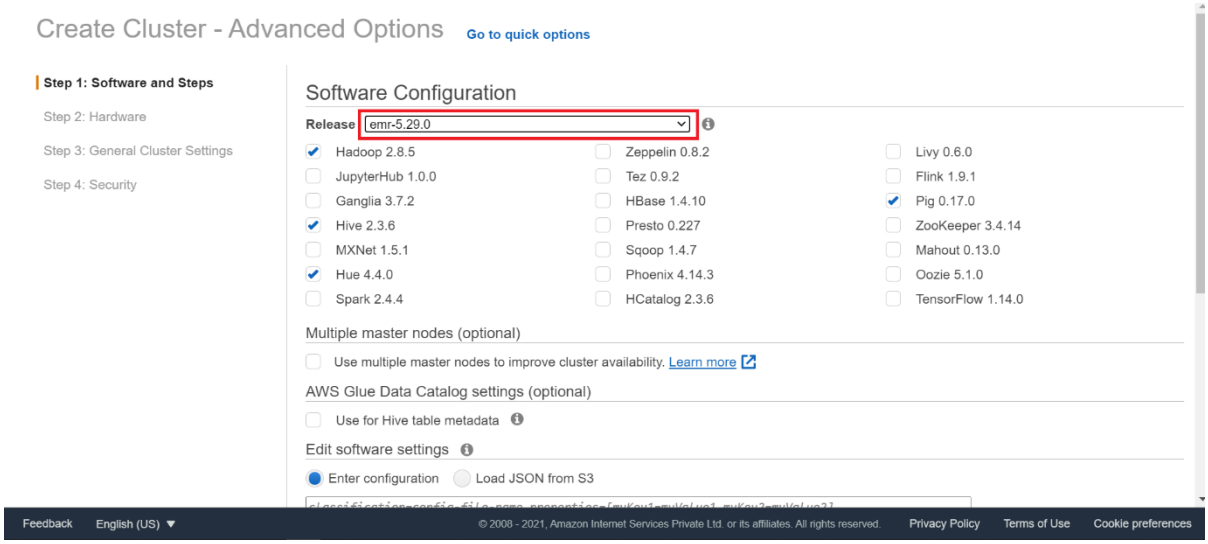

STEP 2: We will be going for a **2-node cluster** for our analysis & we will select **m4.large** instance type each for both master & core node.

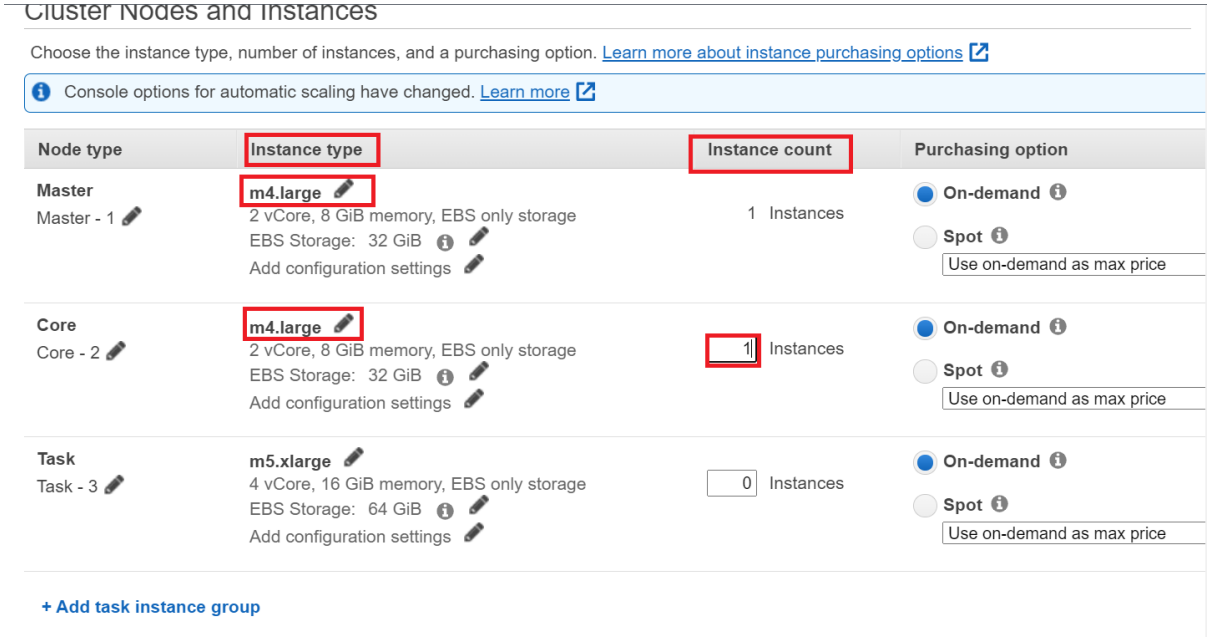

Step 3: Select a cluster name. Here we have taken the cluster name as **My cluster**.

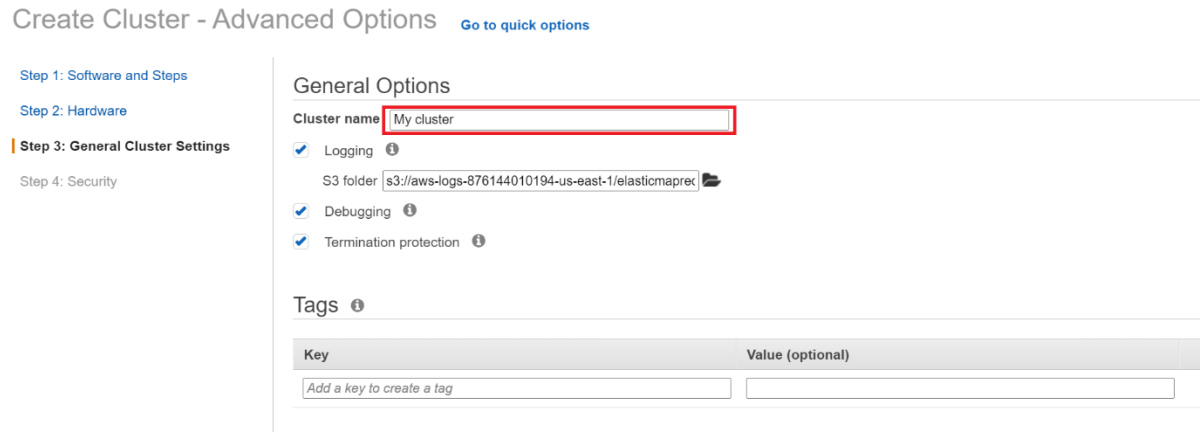

Step 4: Select an already created key-pair which will be used while connecting to master node and **hive_test** is the name of the key that we have used.

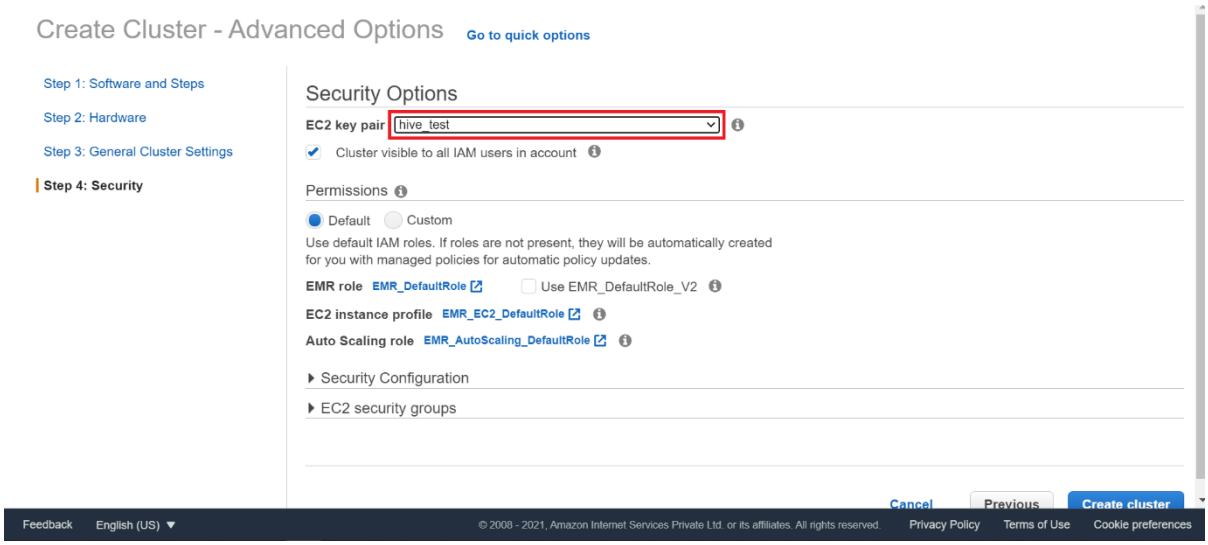

Step 5: Our cluster has been created successfully and is in **waiting state** which indicates its ready to be connected from the local system.

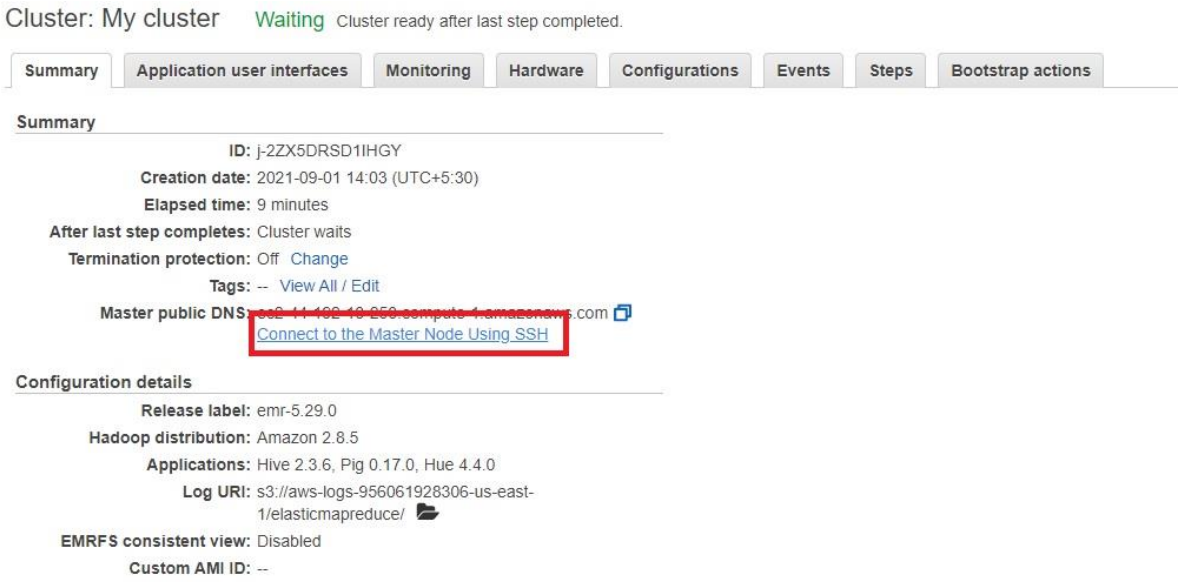

Step 6: Copy the highlighted link i.e., the master public DNS.

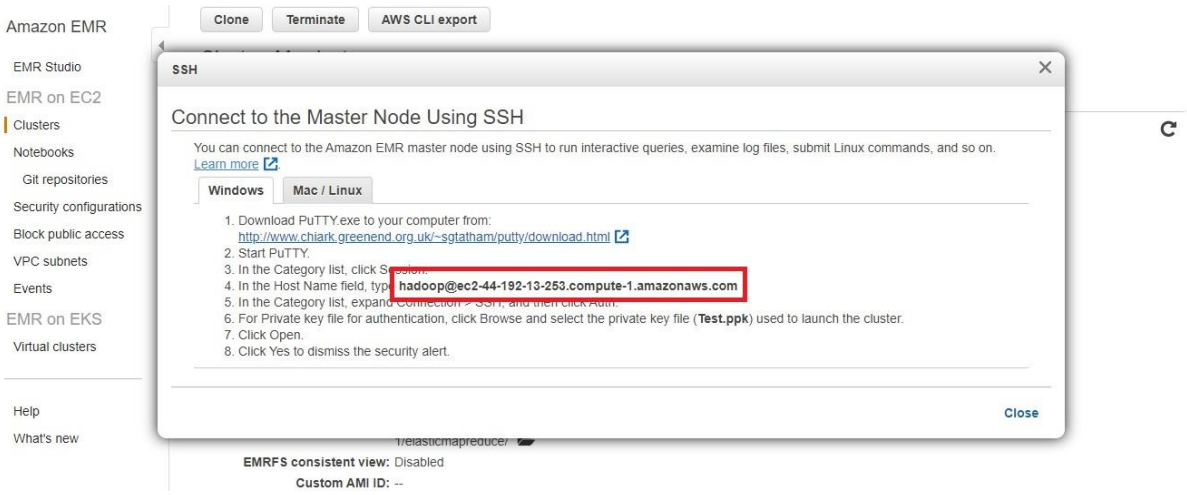

Step 7: Next, paste the address in the **Host Name** field.

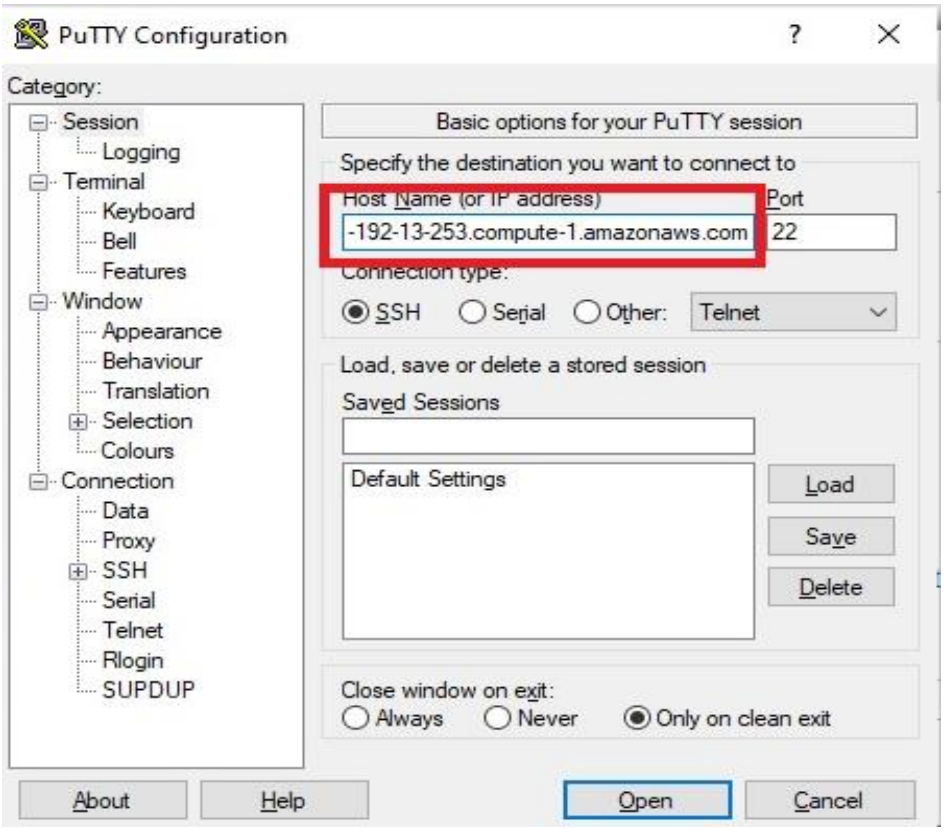

Step 7: Click on **SSH** & then **Auth**. Give the location where the key pair is stored in the local system and click on **Open.**

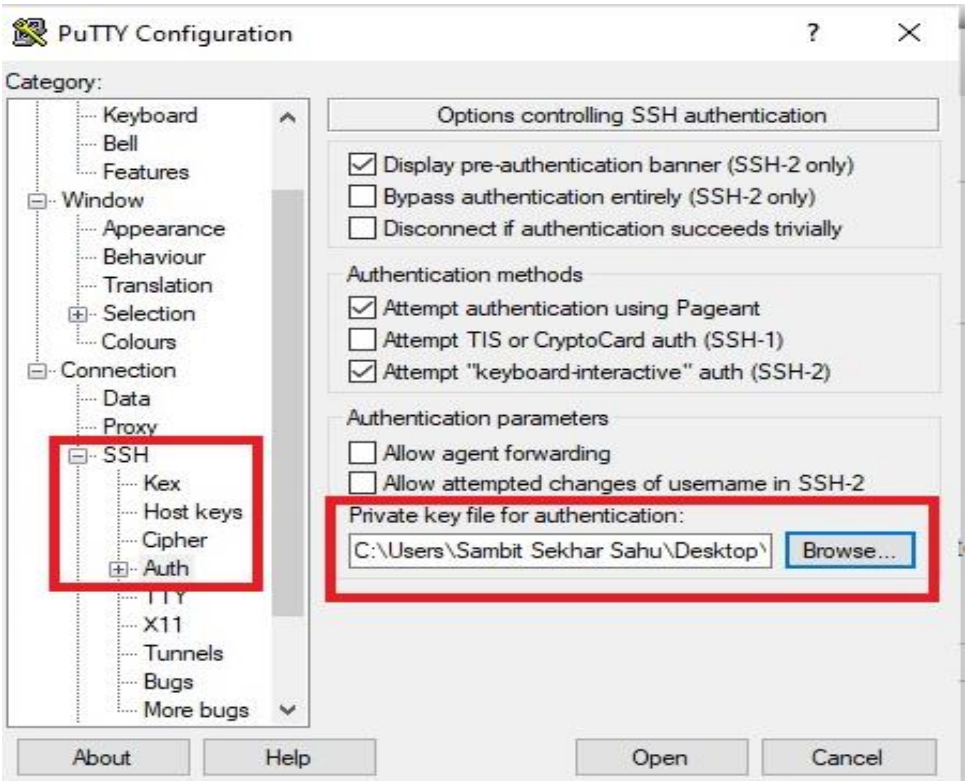

Step 8: Next click on Accept which will open the SSH terminal.

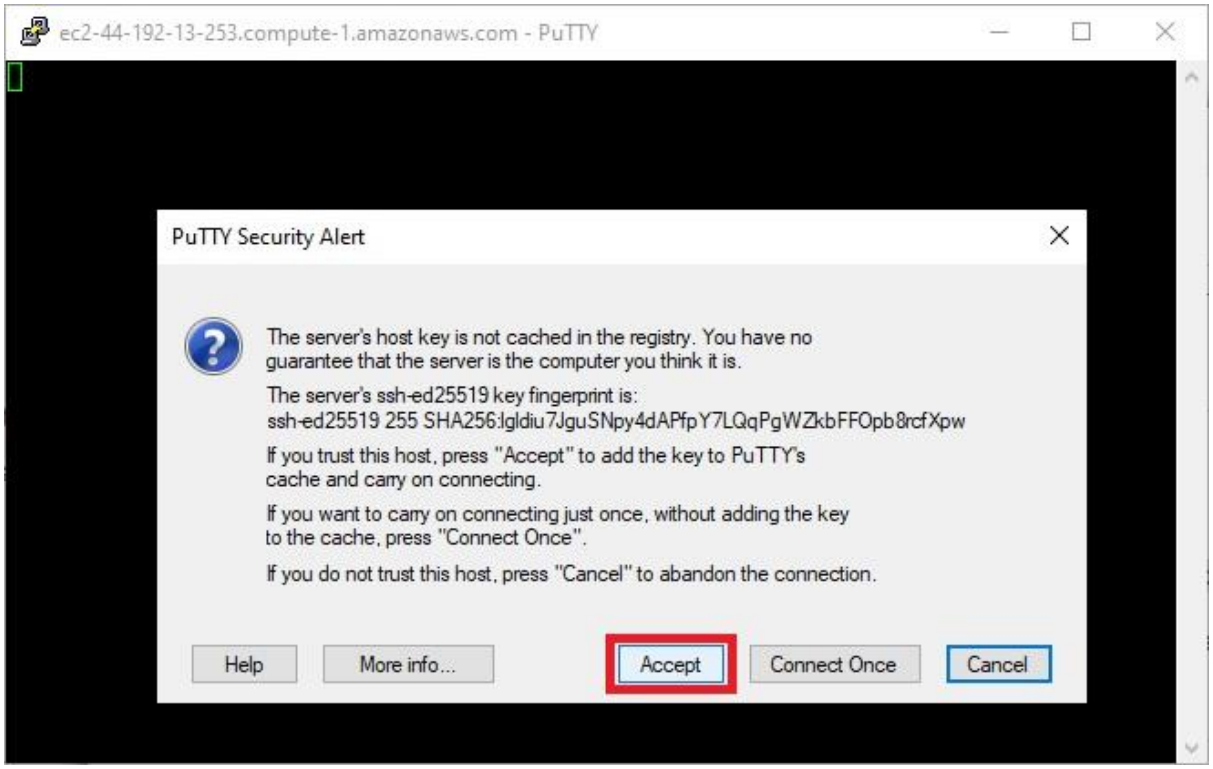

After we have created an EMR cluster & successfully connected to it via putty we can begin to code in the SSH Terminal.

DATA LOADING INTO S3 BUCKET

aws s3 cp s3://e-commerce-events-ml/2019-Oct.csv s3://casestudy-data/2019-oct.csv

aws s3 cp s3://e-commerce-events-ml/2019-Nov.csv s3://casestudy-data/2019-nov.csv

```
[hadoop@ip-172-31-65-119 ~]$ aws s3 cp s3://e-commerce-events-ml/2019-Oct.csv s3://casestudy-data/2019-oct.csv
copy: s3://e-commerce-events-ml/2019-Oct.csv to s3://casestudy-data/2019-oct.csv
[hadoop@ip-172-31-65-119 ~]$ aws s3 cp s3://e-commerce-events-ml/2019-Nov.csv s3://casestudy-data/2019-nov.csv
copy: s3://e-commerce-events-ml/2019-Nov.csv to s3://casestudy-data/2019-nov.csv
```

using the above code, we could directly copy the data files into Hadoop without actually downloading the data files into our system, and the copying completes instantly.

WORKING WITH HDFS

- **Creating a folder in Hadoop**

  hadoop fs -mkdir /cstudy-folder
  hadoop fs -ls /

```
[hadoop@ip-172-31-75-216 ~]$ hadoop fs -mkdir /cstudy-folder
[hadoop@ip-172-31-75-216 ~]$ hadoop fs -ls /
Found 5 items
drwxr-xr-x   - hdfs   hadoop          0 2021-09-01 08:39 /apps
drwxr-xr-x   - hadoop hadoop          0 2021-09-01 08:52 /cstudy-folder
drwxrwxrwt   - hdfs   hadoop          0 2021-09-01 08:41 /tmp
drwxr-xr-x   - hdfs   hadoop          0 2021-09-01 08:39 /user
drwxr-xr-x   - hdfs   hadoop          0 2021-09-01 08:39 /var
[hadoop@ip-172-31-75-216 ~]$
```

- **Copying October & November data from S3 bucket into HDFS**

  hadoop distcp s3://casestudy-data/2019-oct.csv /cstudy-folder/2019-oct.csv
  hadoop distcp s3://casestudy-data/2019-nov.csv /cstudy-folder/2019-nov.csv

```
[hadoop@ip-172-31-75-216 ~]$ hadoop distcp s3://casestudy-data/2019-oct.csv /cst
udy-folder/2019-oct.csv
21/09/01 08:53:30 INFO tools.DistCp: Input Options: DistCpOptions{atomicCommit=f
alse, syncFolder=false, deleteMissing=false, ignoreFailures=false, overwrite=fal
se, skipCRC=false, blocking=true, numListstatusThreads=0, maxMaps=20, mapBandwid
th=100, sslConfigurationFile='null', copyStrategy='uniformsize', preserveStatus=
[], preserveRawXattrs=false, atomicWorkPath=null, logPath=null, sourceFileListin
g=null, sourcePaths=[s3://casestudy-data/2019-oct.csv], targetPath=/cstudy-folde
r/2019-oct.csv, targetPathExists=false, filtersFile='null'}
21/09/01 08:53:30 INFO client.RMProxy: Connecting to ResourceManager at ip-172-3
1-75-216.ec2.internal/172.31.75.216:8032
21/09/01 08:53:34 INFO tools.SimpleCopyListing: Paths (files+dirs) cnt = 1; dirC
nt = 0
21/09/01 08:53:34 INFO tools.SimpleCopyListing: Build file listing completed.
21/09/01 08:53:34 INFO Configuration.deprecation: io.sort.mb is deprecated. Inst
ead, use mapreduce.task.io.sort.mb
21/09/01 08:53:34 INFO Configuration.deprecation: io.sort.factor is deprecated.
Instead, use mapreduce.task.io.sort.factor
21/09/01 08:53:34 WARN hdfs.DataStreamer: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at org.apache.hadoop.hdfs.DataStreamer.closeResponder(DataStreamer.java:
973)
        at org.apache.hadoop.hdfs.DataStreamer.endBlock(DataStreamer.java:624)
        at org.apache.hadoop.hdfs.DataStreamer.run(DataStreamer.java:801)
21/09/01 08:53:34 INFO tools.DistCp: Number of paths in the copy list: 1
21/09/01 08:53:34 INFO tools.DistCp: Number of paths in the copy list: 1
21/09/01 08:53:35 INFO client.RMProxy: Connecting to ResourceManager at ip-172-3
1-75-216.ec2.internal/172.31.75.216:8032
```

```
                S3: Number of large read operations=0
                S3: Number of write operations=0
        Job Counters
                Launched map tasks=1
                Other local map tasks=1
                Total time spent by all maps in occupied slots (ms)=749792
                Total time spent by all reduces in occupied slots (ms)=0
                Total time spent by all map tasks (ms)=23431
                Total vcore-milliseconds taken by all map tasks=23431
                Total megabyte-milliseconds taken by all map tasks=23993344
        Map-Reduce Framework
                Map input records=1
                Map output records=0
                Input split bytes=137
                Spilled Records=0
                Failed Shuffles=0
                Merged Map outputs=0
                GC time elapsed (ms)=287
                CPU time spent (ms)=18480
                Physical memory (bytes) snapshot=567488512
                Virtual memory (bytes) snapshot=3288674304
                Total committed heap usage (bytes)=503316480
        File Input Format Counters
                Bytes Read=217
        File Output Format Counters
                Bytes Written=0
        DistCp Counters
                Bytes Copied=482542278
                Bytes Expected=482542278
                Files Copied=1
```

Similarly, November data file was also copied, as shown in above image.

- **Verifying if data has been copied successfully**

  hadoop fs -ls /cstudy-folder

```
[hadoop@ip-172-31-75-216 ~]$ hadoop fs -ls /cstudy-folder
Found 2 items
-rw-r--r--   1 hadoop hadoop  545839412 2021-09-01 08:57 /cstudy-folder/2019-nov.csv
-rw-r--r--   1 hadoop hadoop  482542278 2021-09-01 08:54 /cstudy-folder/2019-oct.csv
[hadoop@ip-172-31-75-216 ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive>
```

WORKING ON HIVE

- **Getting into Hive CLI – making a new database named *cstudy***

  create database if not exists cstudy ;
  use cstudy ;

```
hive> create database if not exists cstudy ;
OK
Time taken: 0.801 seconds
hive> use cstudy ;
OK
Time taken: 0.055 seconds
hive>
```

- **Creating a common table named clickstream and storing both October & November data in it**

  create external table if not exists clickstream(event_time timestamp, event_type string, product_id string, category_id string, category_code string, brand string, price float, user_id bigint, user_session string) ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' WITH SERDEPROPERTIES ('separatorChar'= ',', 'escapeChar'= '\\') stored as textfile LOCATION 'hdfs:///cstudy-folder/' tblproperties ('skip.header.line.count' = '1') ;

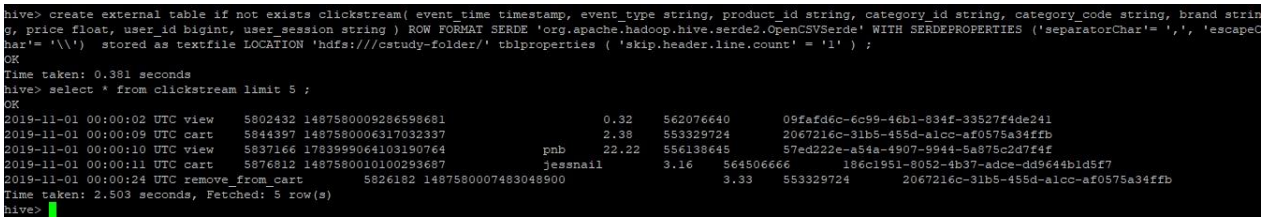

- **To create optimised table having partitions & buckets we need to enable some settings**

  set hive.exec.dynamic.partition = true ;
  set hive.exec.dynamic.partition.mode = nonstrict ;
  set hive.enforce.bucketing = true ;

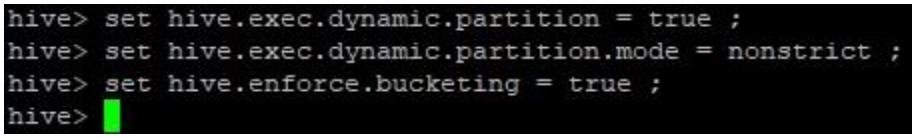

- **Creating table with dynamic partitions and buckets and inserting data into it**

  create table if not exists dynpart_buck_clickstream (event_time string, product_id string, category_id string, category_code string, brand string, price float, user_id bigint, user_session string) partitioned by (event_type string) clustered by (category_code) into 13 buckets row format delimited fields terminated by ',' lines terminated by '\n' stored as textfile;

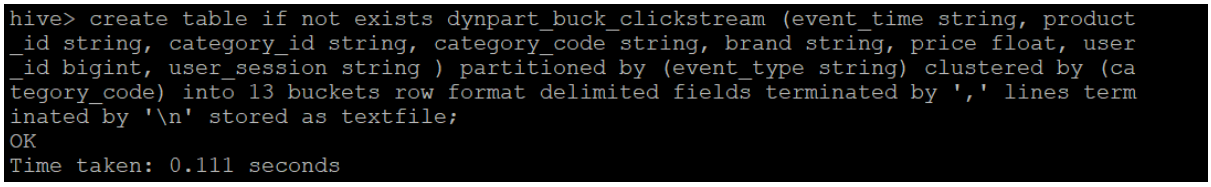

insert into table dynpart_buck_clickstream partition (event_type) select event_time, product_id,category_id,category_code,brand,price,user_id,user_session,event_type from clickstream ;

```
hive> insert into table dynpart_buck_clickstream partition (event_type) select event_time,product_id
,category_id,category_code,brand,price,user_id,user_session,event_type from clickstream ;
Query ID = hadoop_20210905054901_bc1b6f1a-c2ac-46be-8715-05a11b27dfff
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1630819538041_0004)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      5          5        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 160.92 s
--------------------------------------------------------------------------------
Loading data to table cstudy.dynpart_buck_clickstream partition (event_type=null)

Loaded : 4/4 partitions.
        Time taken to load dynamic partitions: 0.789 seconds
        Time taken for adding to write entity : 0.004 seconds
OK
Time taken: 172.112 seconds
```

- **Describing both tables**

describe clickstream;
describe dynpart_buck_clickstream;

```
hive> describe clickstream;
OK
col_name        data_type       comment
event_time              string                  from deserializer
event_type              string                  from deserializer
product_id              string                  from deserializer
category_id             string                  from deserializer
category_code           string                  from deserializer
brand                   string                  from deserializer
price                   string                  from deserializer
user_id                 string                  from deserializer
user_session            string                  from deserializer
Time taken: 0.042 seconds, Fetched: 9 row(s)
hive> describe dynpart_buck_clickstream;
OK
col_name        data_type       comment
event_time              string
product_id              string
category_id             string
category_code           string
brand                   string
price                   float
user_id                 bigint
user_session            string
event_type              string

# Partition Information
# col_name              data_type               comment

event_type              string
Time taken: 0.11 seconds, Fetched: 14 row(s)
```

- **Checking data in both tables**

  set hive.cli.print.header=true;
  *select * from clickstream limit 5 ;*
  *select * from dynpart_buck_clickstream limit 5 ;*

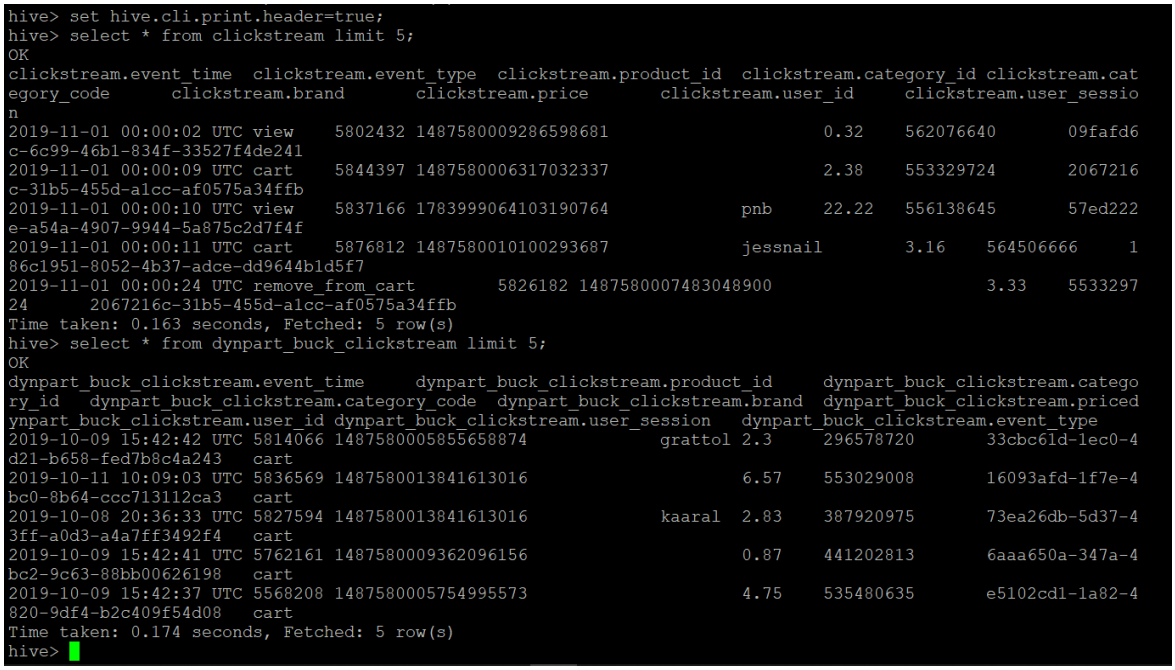

- **Checking if partitions were created successfully**

  show partitions dynpart_buck_clickstream;

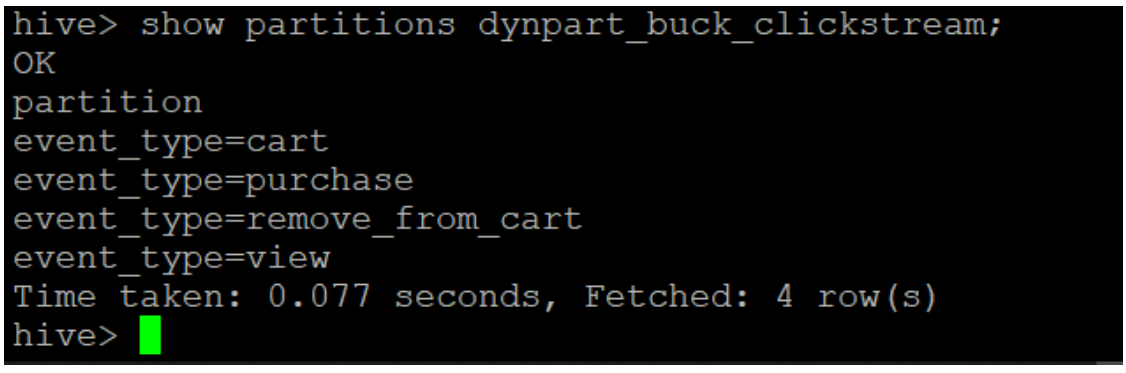

- **Checking in Hadoop if partitions were created successfully**

  hadoop fs -ls /user/hive/warehouse/cstudy.db/dynpart_buck_clickstream

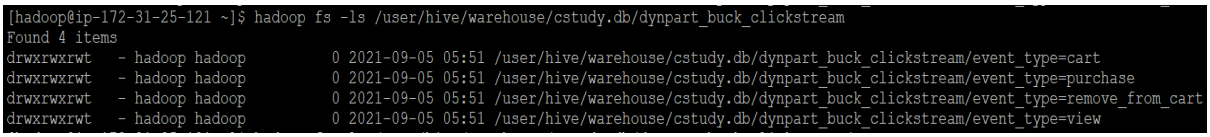

  Overall, we have made two tables,

  - One common table named **clickstream** which contains data of both October & November.
  - One table with partitions & buckets named **dynpart_buck_clickstream** for optimised querying which also contains data of both October & November.

  So, all the preparations are done & now we can move to query analysis-

QUERY ANALYSIS

1. **Find the total revenue generated due to purchases made in October.**

a) *Unoptimized query:*

select sum(price) as total_revenue_oct from clickstream where event_type='purchase' and month(event_time)=10 ;
and month(event_time)=10 ;

```
hive> select sum(price) as total_revenue_oct from clickstream where event_type='purchase' and month(event_time)=10 ;
Query ID = hadoop_20210905123412_c7c81d35-e0aa-460d-8895-d54970d87e71
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630842684204_0005)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 53.59 s
--------------------------------------------------------------------------------
OK
1211538.4299997438
Time taken: 54.549 seconds, Fetched: 1 row(s)
```

b) *Optimized query:*

select sum(price) as total_revenue_oct from dynpart_buck_clickstream where event_type='purchase' and month(event_time)=10 ;

```
hive> select sum(price) as total_revenue_oct from dynpart_buck_clickstream where event_type='purchase' and month(event_time)=10 ;
Query ID = hadoop_20210905084209_249c7677-7a15-428e-a5ee-7b0243706379
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630828819874_0004)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 15.03 s
--------------------------------------------------------------------------------
OK
1211538.4295325726
Time taken: 15.764 seconds, Fetched: 1 row(s)
```

Optimized query using dynpart_buck_clickstream table

*The total revenue generated in October is **1211538.429.** Optimized query took **15.764 secs** while unoptimized query took **54.549 secs** to fetch the same result.*

2. Write a query to yield the total sum of purchases per month in a single output.

a) *<u>Unoptimized query:</u>*
select month(event_time) as month, sum(price) as total_revenue from clickstream where event_type='purchase' group by month(event_time);

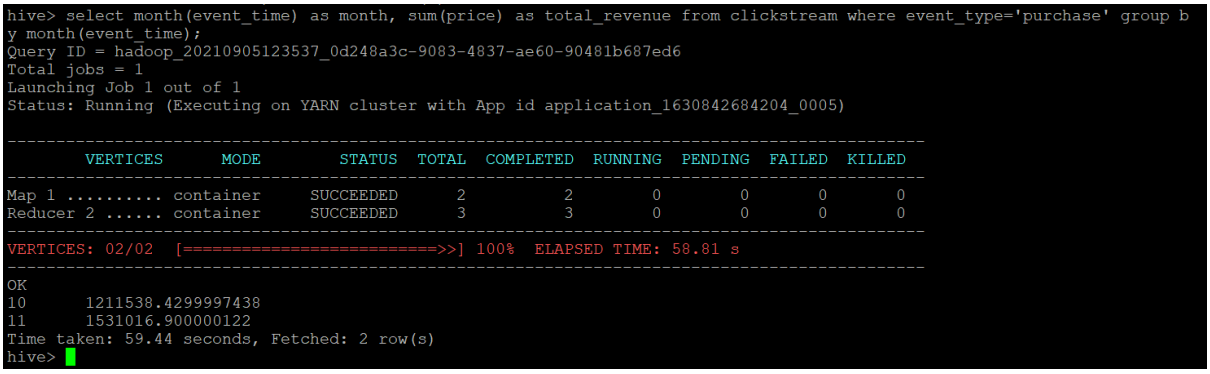

b) *<u>Optimized query:</u>*
select month(event_time) as month, sum(price) as total_revenue from dynpart_buck_clickstream where event_type='purchase' group by month(event_time);

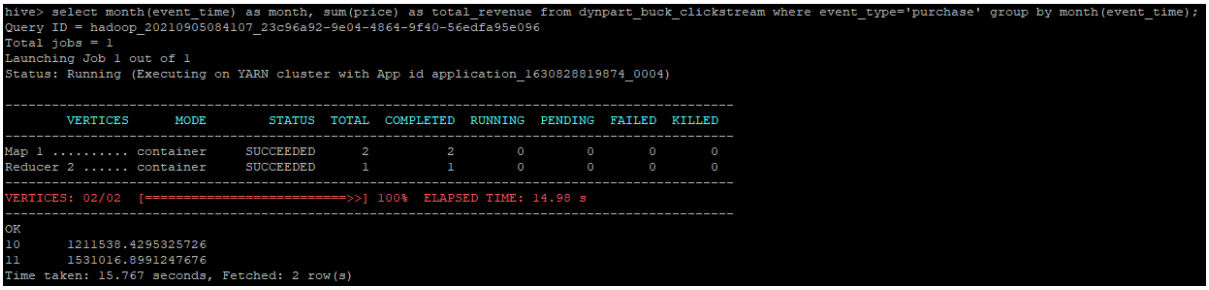

Optimized query using dynpart_buck_clickstream table

*Total sum of purchases for **October** is **1211538.429** while for the **November** it's **1531016.899.** Optimized query took **15.767 secs** while unoptimized query took **59.44 secs**.*

3. **Write a query to find the change in revenue due to purchases from October to November.**

a) *<u>Unoptimized query:</u>*
select (sum(case when month(event_time)=11 then price else 0 end) - sum(case when month(event_time)=10 then price else 0 end)) as change_in_revenue from clickstream where event_type='purchase' ;

```
hive> select (sum(case when month(event_time)=11 then price else 0 end) - sum(case when month(event_time)=10 then price e
lse 0 end)) as change_in_revenue from clickstream where event_type='purchase' ;
Query ID = hadoop_20210905123723_786786e1-2fd2-45b7-8ac7-521884bc81a6
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630842684204_0005)

----------------------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
----------------------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
----------------------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 55.58 s
----------------------------------------------------------------------------------------------
OK
319478.4700003781
Time taken: 56.121 seconds, Fetched: 1 row(s)
hive>
```

b) *<u>Optimized query:</u>*
select (sum(case when month(event_time)=11 then price else 0 end) - sum(case when month(event_time)=10 then price else 0 end)) as change_in_revenue from dynpart_buck_clickstream where event_type='purchase' ;

```
hive> select (sum(case when month(event_time)=11 then price else 0 end) - sum(case when month(event_time)=10 then price else 0 end)) as change_in_revenue from dynpart_b
uck_clickstream where event_type='purchase' ;
Query ID = hadoop_20210905084527_511cf0e7-a0ae-42df-858d-25a4f88930fa
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630828819874_0004)

----------------------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
----------------------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
----------------------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 15.17 s
----------------------------------------------------------------------------------------------
OK
319478.469592195
Time taken: 15.964 seconds, Fetched: 1 row(s)
hive>
```

Optimized query using dynpart_buck_clickstream table

*Change in revenue is **319478.469**. Optimized query took **15.964 secs** while unoptimized query took **56.121 secs**.*

4. **Find distinct categories of products. Categories with null category code can be ignored.**

a) *<u>Unoptimized query:</u>*
select distinct(category_code) from clickstream where category_code != '' ;

```
hive> select distinct(category_code) from clickstream where category_code != '' ;
Query ID = hadoop_20210904062628_11ce72fd-060f-413e-8dba-84b509169cdb
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630735375373_0004)

----------------------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
----------------------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      5          5        0        0       0       0
----------------------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 56.41 s
----------------------------------------------------------------------------------------------
OK
accessories.cosmetic_bag
stationery.cartrige
accessories.bag
appliances.environment.vacuum
furniture.living_room.chair
sport.diving
appliances.personal.hair_cutter
appliances.environment.air_conditioner
apparel.glove
furniture.bathroom.bath
furniture.living_room.cabinet
Time taken: 57.076 seconds, Fetched: 11 row(s)
```

b) *<u>Optimized query:</u>*
select distinct(category_code) from dynpart_buck_clickstream where category_code != '' ;

```
hive> select distinct(category_code) from dynpart_buck_clickstream where category_code != '' ;
Query ID = hadoop_20210905084623_a36e53ed-c006-4b97-bf65-237ed6c9533c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630828819874_0004)

----------------------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
----------------------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      7          7        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      4          4        0        0       0       0
----------------------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 27.92 s
----------------------------------------------------------------------------------------------
OK
accessories.bag
appliances.environment.vacuum
appliances.personal.hair_cutter
sport.diving
apparel.glove
furniture.bathroom.bath
furniture.living_room.cabinet
stationery.cartrige
accessories.cosmetic_bag
appliances.environment.air_conditioner
furniture.living_room.chair
Time taken: 28.596 seconds, Fetched: 11 row(s)
```

Optimized query using dynpart_buck_clickstream table

*There are **11 distinct categories**. Optimized query took **28.596 secs** while unoptimized query took **57.076 secs**.*

5. Find the total number of products available under each category.

a) *<u>Unoptimized query:</u>*

select category_code, count(product_id) as total_products from clickstream where category_code != '' group by category_code ;

```
hive> select category_code, count(product_id) as total_products from clickstream where category_code != ''
 group by category_code ;
Query ID = hadoop_20210904063003_79ff5b37-dd18-4f4f-ac22-3004c25efd23
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630735375373_0004)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      5          5        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 54.80 s
--------------------------------------------------------------------------------
OK
accessories.cosmetic_bag        1248
stationery.cartrige     26722
accessories.bag 11681
appliances.environment.vacuum   59761
furniture.living_room.chair     308
sport.diving    2
appliances.personal.hair_cutter 1643
appliances.environment.air_conditioner  332
apparel.glove   18232
furniture.bathroom.bath 9857
furniture.living_room.cabinet   13439
Time taken: 55.414 seconds, Fetched: 11 row(s)
```

b) *<u>Optimized query:</u>*

select category_code, count(product_id) as total_products from dynpart_buck_clickstream where category_code != '' group by category_code ;

```
hive> select category_code, count(product_id) as total_products from dynpart_buck_clickstream where category_code != '' group by category_code ;
Query ID = hadoop_20210905084842_8d776c68-7b05-4920-a49a-457b8b74f8af
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630828819874_0004)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      7          7        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      4          4        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 26.65 s
--------------------------------------------------------------------------------
OK
accessories.bag 11681
appliances.environment.vacuum   59761
appliances.personal.hair_cutter 1643
sport.diving    2
apparel.glove   18232
furniture.bathroom.bath 9857
furniture.living_room.cabinet   13439
stationery.cartrige     26722
accessories.cosmetic_bag        1248
appliances.environment.air_conditioner  332
furniture.living_room.chair     308
Time taken: 27.253 seconds, Fetched: 11 row(s)
hive>
```

Optimized query using dynpart_buck_clickstream table

*Optimized query took **27.253 secs** while unoptimized query took **55.414 secs**.*

6. Which brand had the maximum sales in October and November combined?

a) *Unoptimized query:*

select brand, sum(price) as total_sales from clickstream where event_type='purchase' and brand != '' group by brand order by total_sales desc limit 1 ;

```
hive> select brand, sum(price) as total_sales from clickstream where event_type='purchase' and brand != '' group by brand
 order by total_sales desc limit 1 ;
Query ID = hadoop_20210905123110_91594964-504b-4671-8547-29359b95ac88
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1630842684204_0005)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      3          3        0        0       0       0
Reducer 3 ...... container     SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 03/03  [==========================>>] 100%  ELAPSED TIME: 58.47 s
--------------------------------------------------------------------------------
OK
runail  148297.9400000003
Time taken: 67.809 seconds, Fetched: 1 row(s)
```

b) *Optimized query:*

select brand, sum(price) as total_sales from dynpart_buck_clickstream where event_type='purchase' and brand != '' group by brand order by total_sales desc limit 1 ;

```
hive> select brand, sum(price) as total_sales from dynpart_buck_clickstream where event_type='purchase' and brand != '' group by brand order by total_sales d
esc limit 1 ;
Query ID = hadoop_20210904084445_0e57640c-c775-4e1f-a491-d81fa5f9f2e6
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630735375373_0011)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
Reducer 3 ...... container     SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 03/03  [==========================>>] 100%  ELAPSED TIME: 16.28 s
--------------------------------------------------------------------------------
OK
runail  148297.93999999692
Time taken: 16.98 seconds, Fetched: 1 row(s)
hive>
```

Optimized query using dynpart_buck_clickstream table

*The top brand is **runail** with **total_sales 148297.939** & to show this optimized query took **16.98 secs** while unoptimized query took **67.809 secs.***

7. Which brands increases their sales from October to November?

a) *<u>Unoptimized query:</u>*

with brand_sales_summary as (select brand, sum(case when month(event_time)=10 then price else 0 end) as oct_sales, sum(case when month(event_time)=11 then price else 0 end) as nov_sales from clickstream where event_type='purchase' group by brand) select brand from brand_sales_summary where (nov_sales-oct_sales)>0 ;

```
hive> with brand_sales_summary as ( select brand, sum(case when month(event_time)=10 then price else 0 end) as oct_sales,
 sum(case when month(event_time)=11 then price else 0 end) as nov_sales from clickstream where event_type='purchase' grou
p by brand ) select brand from brand_sales_summary where (nov_sales-oct_sales)>0 ;
Query ID = hadoop_20210905120034_838a5924-9fdd-40d3-a3cc-869d82467c70
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630842684204_0003)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 ..... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ..... container     SUCCEEDED      3          3        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 61.58 s
--------------------------------------------------------------------------------
OK
airnails
artex
binacil
bioaqua
blixz
bluesky
bpw.style
carmex
chi
concept
cosima
cosmoprofi
deoproce
depilflax
dewal
dizao
egomania
elizavecca
ellips
finish
freshbubble
grattol
haruyama
helloganic
insight
```

```
beauty-free
beautyblender
benovy
candy
coifin
cristalinas
cutrin
domix
ecocraft
elskin
enjoy
entity
eos
estel
estelare
farmavita
fedua
foamie
glysolid
godefroy
inm
irisk
kamill
kares
kaypro
keen
kinetics
koelcia
lianail
lowence
matreshka
mavala
missha
moyou
nagaraku
profepil
rasyan
refectocil
skinity
smart
solomeya
swarovski
trind
uno
yu-r
Time taken: 62.784 seconds, Fetched: 161 row(s)
```

b) *<u>Optimized query:</u>*

with brand_sales_summary as (select brand, sum(case when month(event_time)=10 then price else 0 end) as oct_sales, sum(case when month(event_time)=11 then price else 0 end) as nov_sales from dynpart_buck_clickstream where event_type='purchase' group by brand) select brand from brand_sales_summary where (nov_sales-oct_sales)>0 ;

```
hive> with brand_sales_summary as ( select brand, sum(case when month(event_time)=10 then price else 0 end) as oct_sales,
 sum(case when month(event_time)=11 then price else 0 end) as nov_sales from dynpart_buck_clickstream where event_type='p
urchase' group by brand ) select brand from brand_sales_summary where (nov_sales-oct_sales)>0 ;
Query ID = hadoop_20210905120229_471b103a-0a54-43b9-afd0-618b06bb8a69
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630842684204_0003)

----------------------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
----------------------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
----------------------------------------------------------------------------------------------
VERTICES: 02/02  [==========================>>] 100%  ELAPSED TIME: 16.00 s
----------------------------------------------------------------------------------------------
OK

airnails
art-visage
artex
aura
balbcare
barbie
batiste
beautix
beauty-free
beautyblender
beauugreen
benovy
binacil
bioaqua
biore
blixz
bluesky
bodyton
bpw.style
browxenna
candy
carmex
chi
coifin
```

```
polarus
profepil
profhenna
protokeratin
provoc
rasyan
refectocil
rosi
roubloff
runail
s.care
sanoto
severina
shary
shik
skinity
skinlite
smart
soleo
solomeya
sophin
staleks
strong
supertan
swarovski
tertio
treaclemoon
trind
uno
uskusi
veraclara
vilenta
yoko
yu-r
zeitun
Time taken: 17.222 seconds, Fetched: 161 row(s)
```

*There is a total of **161 brands** that had increased sales from October to November. To show this optimized query took **17.222 secs** while unoptimized query took **62.784 secs.***

8. **Your company wants to reward the top 10 users of its website with a Golden Customer plan. Write a query to generate a list of top 10 users who spend the most.**

a) *<u>Unoptimized query:</u>*
with spending_summ as (select user_id, sum(price) as overall_spending from clickstream where event_type='purchase' group by user_id order by overall_spending desc) select user_id from spending_summ limit 10 ;

```
hive> with spending_summ as (select user_id, sum(price) as overall_spending from clickstream where event_type='purchase'
group by user_id order by overall_spending desc) select user_id from spending_summ limit 10 ;
Query ID = hadoop_20210905120758_c8d6d0bd-a576-4fad-b314-8eb07c66ba9a
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630842684204_0003)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 .......... container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      3          3        0        0       0       0
Reducer 3 ...... container     SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 03/03  [==========================>>] 100%  ELAPSED TIME: 61.04 s
--------------------------------------------------------------------------------
OK
557790271
150318419
562167663
531900924
557850743
522130011
561592095
431950134
566576008
521347209
Time taken: 61.985 seconds, Fetched: 10 row(s)
hive>
```

b) *Optimized query:*

with spending_summ as (select user_id, sum(price) as overall_spending from dynpart_buck_clickstream where event_type='purchase' group by user_id order by overall_spending desc) select user_id from spending_summ limit 10 ;

```
hive> with spending_summ as (select user_id, sum(price) as overall_spending from dynpart_buck_clickstream where event_typ
e='purchase' group by user_id order by overall_spending desc) select user_id from spending_summ limit 10 ;
Query ID = hadoop_20210905120657_3b17e114-263d-4e8e-91f9-dda549d341f3
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1630842684204_0003)

----------------------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
----------------------------------------------------------------------------------------------
Map 1 ..........   container     SUCCEEDED      2          2        0        0       0       0
Reducer 2 ......   container     SUCCEEDED      1          1        0        0       0       0
Reducer 3 ......   container     SUCCEEDED      1          1        0        0       0       0
----------------------------------------------------------------------------------------------
VERTICES: 03/03  [==========================>>] 100%  ELAPSED TIME: 14.79 s
----------------------------------------------------------------------------------------------
OK
557790271
150318419
562167663
531900924
557850743
522130011
561592095
431950134
566576008
521347209
Time taken: 15.688 seconds, Fetched: 10 row(s)
hive>
```

Optimized query using dynpart_buck_clickstream table

*So above are the top 10 customers who spend the most and should be awarded with golden customer plan. To show this the optimized query took **15.688 secs** while unoptimized query took **61.985 secs**.*

- **Dropping database**

drop database cstudy cascade;

```
hive> drop database cstudy cascade;
OK
Time taken: 0.33 seconds
```

show databases;

```
hive> show databases;
OK
default
Time taken: 0.039 seconds, Fetched: 1 row(s)
```

We are done with our analysis. Finally, we will move towards terminating our cluster

TERMINATING THE EMR CLUSTER

Step 1: Click on Terminate.

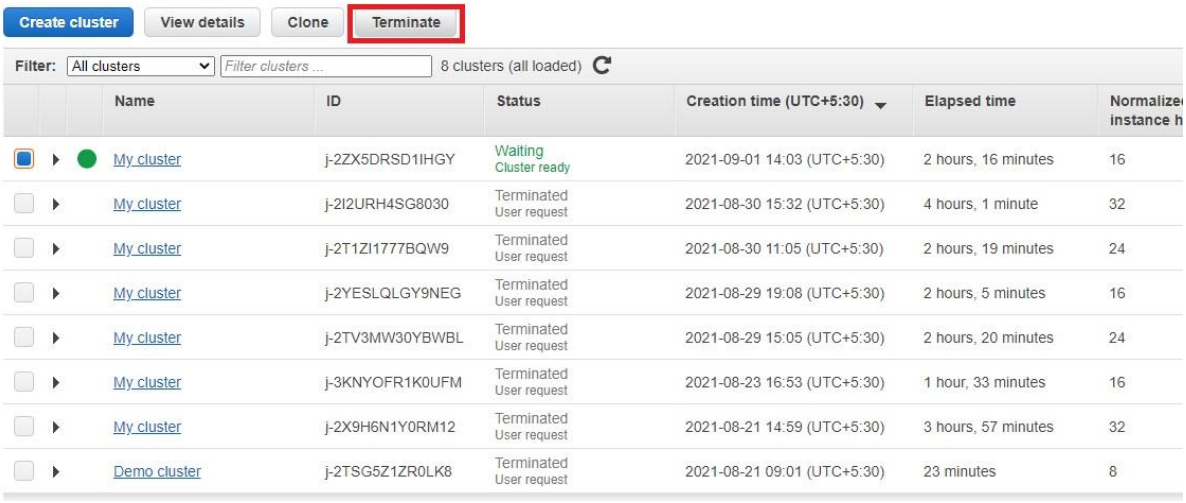

Step 2: A security/confirmation message will pop up. Click on Terminate.

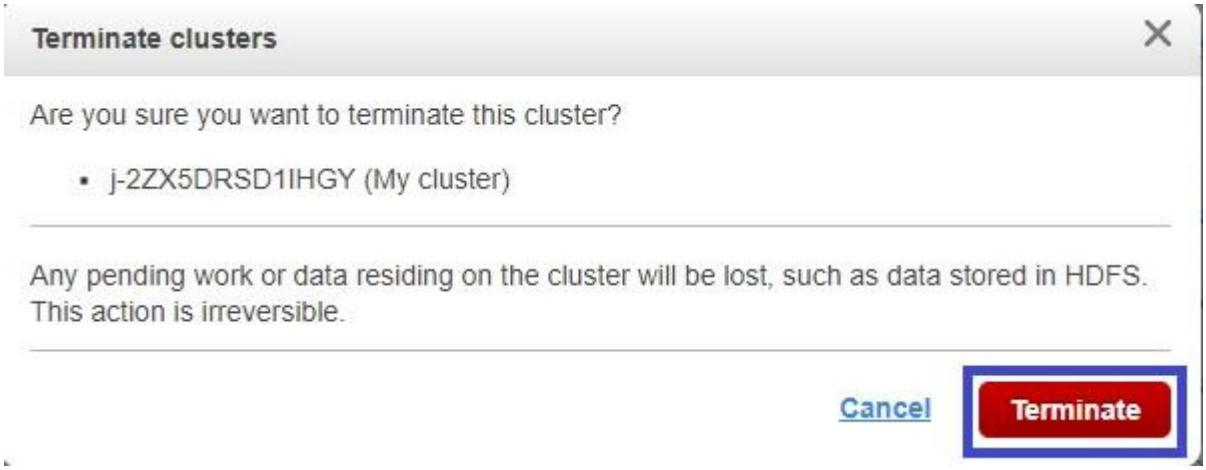

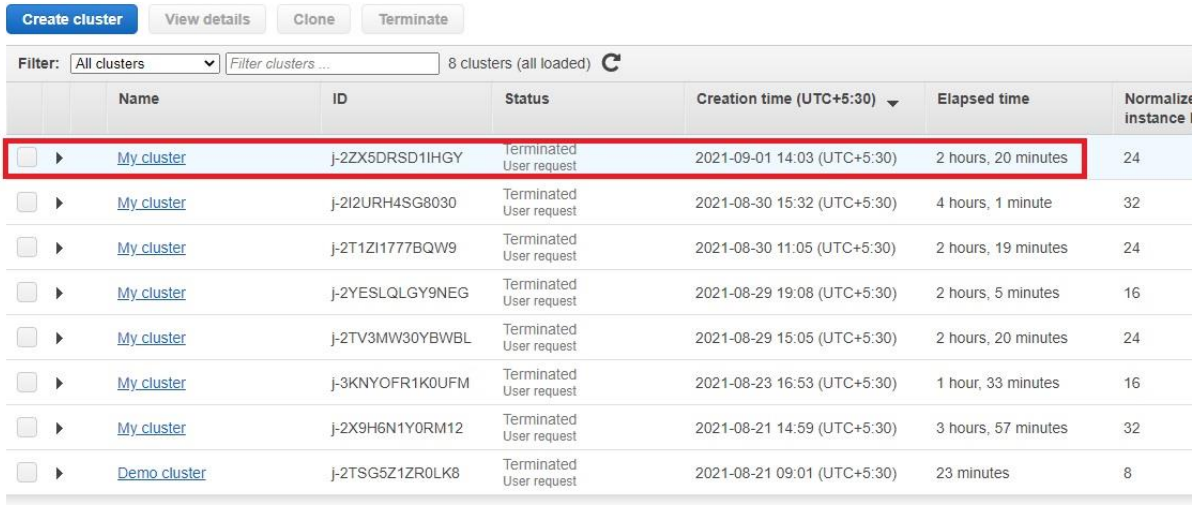

The Cluster has been terminated. Now we can log out of our AWS account.